

# CS426: Hadoop Experiment Report

Zuoming Zhang  
5120309626  
237632270@qq.com  
Department of Computer  
Science and Engineering  
School of Electronic Information  
and Electrical Engineering  
Shanghai Jiao Tong University

Weichen Li  
5120309662  
eizo.lee@sjtu.edu.cn  
Department of Computer  
Science and Engineering  
School of Electronic Information  
and Electrical Engineering  
Shanghai Jiao Tong University

Peiyuan Liu  
5120309394  
ms\_green@sjtu.edu.cn  
Department of Computer  
Science and Engineering  
School of Electronic Information  
and Electrical Engineering  
Shanghai Jiao Tong University

## Abstract

This report describes steps to construct Hadoop both in **Single Mode** and in **Cluster Mode**. Our report & hadoop installation guide is suitable for both real computers and virtual machines, and the operating system is based on **Ubuntu** (both Ubuntu Desktop and Ubuntu Server are OK, we highly recommend you to use Ubuntu Desktop for Master and Server for Slaves). You can follow our report step by step and build your own hadoop in single/cluster mode. We also run *WordCount* program, in both **Single Mode** and **Cluster Mode**, to test the performance of our hadoop system. Then we further discussed our understanding of hadoop and some interesting findings in the conclusion part.

## CONTENTS

<b>I</b>	<b>Division of Labor</b>	2
<b>II</b>	<b>Experiment Environment</b>	2
II-A	Network Environment . . . . .	2
II-B	Operating System Environment . . . . .	2
<b>III</b>	<b>Single Mode Construction</b>	2
III-A	Add the Hadoop user . . . . .	2
III-B	Add authority to the Hadoop user . . . . .	3
III-C	Login Ubuntu using the new Hadoop user . . . . .	3
III-D	Installation of ssh . . . . .	3
III-E	Installation of Java environment . . . . .	5
III-F	Installation of Hadoop 2.6.0 . . . . .	5
<b>IV</b>	<b>WordCount Test in Single Mode</b>	6
<b>V</b>	<b>Cluster Mode Construction</b>	7
V-A	Preparations for constructing cluster mode. . . . .	7
V-B	SSH password-free login on slaves . . . . .	8
V-C	Cluster environment configuration . . . . .	9
<b>VI</b>	<b>WordCount Test in Cluster Mode</b>	10
VI-A	Format HDFS . . . . .	10
VI-B	Start HDFS . . . . .	10
VI-C	Execute WordCount . . . . .	11
VI-D	Stop HDFS . . . . .	13
<b>VII</b>	<b>Conclusion &amp; Additional Finding</b>	13
VII-A	Conclusion . . . . .	13
VII-B	Additional Finding . . . . .	16
	<b>References</b>	16

## I. DIVISION OF LABOR

Our group is composed of three members: Zuoming Zhang, Weichen Li, Peiyuan Liu, and Zuoming Zhang is the team leader.

In the first project, all the three of us members have implemented the single mode separately. To construct the cluster mode and implement this report & hadoop installation guide, we divide the labor in the following pattern according to each one's ability and experience:

When first trying to establish the cluster mode, Weichen Li gathered relevant information and guidance from the Internet. Then his computer worked as the Master node to control the procedure of construction. Most of the configuration and settings in the cluster mode are implemented by him.

After the first success, Zuoming Zhang reimplemented the whole process in standard steps without mistakes, get the screenshots and write a rough Chinese version draft. He is also responsible for solving technique problems based on his hadoop experience in pseudo-cluster mode.

Then Peiyuan Liu further implemented the project report in English in detail based on the Chinese draft, and she is responsible for setting up the type of the whole report.

## II. EXPERIMENT ENVIRONMENT

### A. Network Environment

In our experiment, we use **real Ethernet** connecting three different laptop computers to build our Hadoop cluster. Each operating system is running on a virtual machine on separate laptops.

We use **VMWARE** to run our operating systems, before we try to connect to the real Ethernet, we must modify the **Network Adapter** of the virtual machine first: Changing **NAT Mode** to **Bridge Mode** (Fig.1).

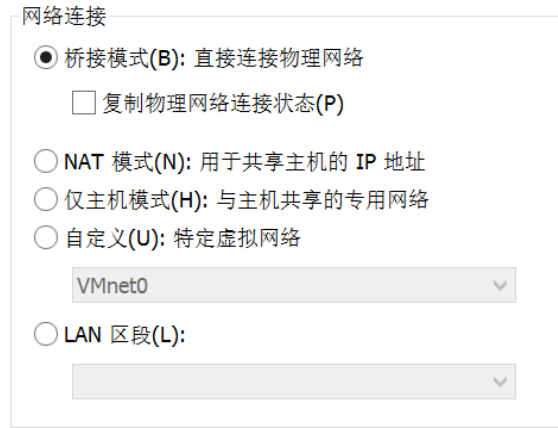


Fig. 1. Modify Network Adapter

### B. Operating System Environment

The operating systems and there kernel versions are listed in Table.I:

TABLE I  
OPERATING SYSTEM ENVIRONMENT

NodeName	Operating System	Kernel Version
Master	Ubuntu Desktop 14.04 64bit	3.13.0-39
Slave1	Ubuntu Kylin 14.10 64bit	3.16.0-23
Slave2	Ubuntu Kylin 14.10 64bit	3.16.0-23

## III. SINGLE MODE CONSTRUCTION

### A. Add the Hadoop user

First we add the **hadoop** user group:

```
sudo addgroup hadoop
```

Then we add the **hadoop** user:

```
sudo adduser --ingroup hadoop hadoop
```

When you type in the "Enter" key, it will ask you to type in the new UNIX password and your personal information. The password is compulsory while you can leave personal information blank.

#### B. Add authority to the Hadoop user

```
sudo gedit /etc/sudoers
```

This command give user **hadoop** the same authority as **root**.

In gedit, add such a line as in Figure 2.

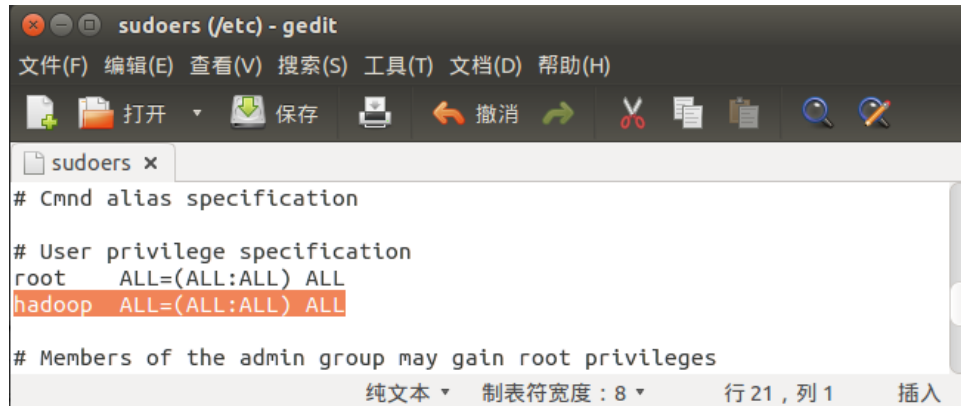


Fig. 2. Add authority to user hadoop

#### C. Login Ubuntu using the new Hadoop user

#### D. Installation of ssh

```
sudo apt-get install openssh-server
```

After the installation, check whether the service is established.

```
ps -e | grep ssh
```

A correct response should look like the codes in Figure 3.

```
hadoop@Master:~$ ps -e | grep ssh
6434 ?        00:00:00 sshd
```

Fig. 3. SSH Service Establishment Check

Then set the password-free login, we will get the public key and the private key.

```
ssh-keygen -t rsa -P ""
```

Corresponding output should look like Figure 4.

At the same time, there will be two new files in directory ~/.ssh, as you can see in Figure 5.

Add the public key to a new file *authorized\_keys*:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

Then we will get files in Figure 6.

Next, login ssh:

```
ssh localhost
```

Result is in Figure 7.

Then exit ssh.

```
exit
```

```

hadoop@Master:~$ ps -e | grep ssh
6434 ?          00:00:00 sshd
hadoop@Master:~$ ssh-keygen -t rsa -P ""
Generating public/private rsa key pair.
Enter file in which to save the key (/home/hadoop/.ssh/id_rsa):
Created directory '/home/hadoop/.ssh'.
Your identification has been saved in /home/hadoop/.ssh/id_rsa.
Your public key has been saved in /home/hadoop/.ssh/id_rsa.pub.
The key fingerprint is:
49:d9:58:66:77:99:b1:04:77:f6:ad:15:2c:7b:20:e0 hadoop@Master
The key's randomart image is:
+---[ RSA 2048]-----+
|          . =  o . * = o |
|          . B o  * o = + |
|          + E . . = + |
|          . . . . + |
|          S      o |
|                    |
|                    |
|                    |
+-----+

```

Fig. 4. Generating SSH Keys

```

hadoop@Master:~$ ls ~/.ssh
id_rsa  id_rsa.pub

```

Fig. 5. SSH files after generating keys

```

hadoop@Master:~$ ls ~/.ssh
authorized_keys  id_rsa  id_rsa.pub

```

Fig. 6. SSH files after adding *authorized\_keys*

```

hadoop@Master:~$ ssh localhost
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is 96:6b:9f:27:4f:67:c6:ee:7a:63:59:4b:1d:86:f2:4b.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'localhost' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 14.10 (GNU/Linux 3.16.0-23-generic x86_64)

* Documentation:  https://help.ubuntu.com/

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

```

Fig. 7. Result of ssh to localhost

### E. Installation of Java environment

```
sudo apt-get install openjdk-7-jdk
```

Check the java environment installation(Fig.8):

```
java -version
```

```
hadoop@Master:~$ java -version
java version "1.7.0_75"
OpenJDK Runtime Environment (IcedTea 2.5.4) (7u75-2.5.4-1~utopic1)
OpenJDK 64-Bit Server VM (build 24.75-b04, mixed mode)
```

Fig. 8. Check the installation of java

### F. Installation of Hadoop 2.6.0

Download Hadoop 2.6.0 from the website <http://mirror.bit.edu.cn/apache/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>  
Unzip and move the hadoop file:

```
sudo tar xzf hadoop-2.6.0.tar.gz
sudo mv hadoop-2.6.0 /usr/local/hadoop
```

Change the authority of the hadoop folder:

```
sudo chmod 777 /usr/local/hadoop
```

Check the location of java(see in Figure 9):

```
update-alternatives --config java
```

```
hadoop@Master:~$ update-alternatives --config java
链接组 java (提供 /usr/bin/java)中只有一个候选项：/usr/lib/jvm/java-7-openjdk-amd64/jre/bin/java
无需配置。
```

Fig. 9. Check the location of java

Here, the java location is: */usr/lib/jvm/java-7-openjdk-amd64*.

Then set *~/bashrc*:

```
sudo gedit ~/.bashrc
```

Add these content to the end of the file:

```
#HADOOP VARIABLES START
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64 #Same to the java location
export HADOOP_INSTALL=/usr/local/hadoop
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"
#HADOOP VARIABLES END
```

Enable the environment variable:

```
source ~/.bashrc
```

Modify *hadoop-env.sh*:

```
sudo gedit /usr/local/hadoop/etc/hadoop/hadoop-env.sh
```

Find **JAVA\_HOME**, modify it as:

```
# The java implementation to use.  
export JAVA_HOME=/usr/lib/jvm/java-7-openjdk-amd64
```

Now, we can run hadoop in **Single Mode**.

#### IV. WORDCOUNT TEST IN SINGLE MODE

```
cd /usr/local/hadoop/  
mkdir input  
cp etc/hadoop/*.xml input  
bin/hadoop jar share/hadoop/mapreduce/sources/hadoop-mapreduce-examples-2.6.0-sources.jar org.apache.hadoop.  
examples.WordCount input output
```

The execution result is in Figure 10

```
15/03/11 00:42:13 INFO mapreduce.Job: Counters: 33  
File System Counters  
  FILE: Number of bytes read=2702897  
  FILE: Number of bytes written=4889838  
  FILE: Number of read operations=0  
  FILE: Number of large read operations=0  
  FILE: Number of write operations=0  
Map-Reduce Framework  
  Map input records=745  
  Map output records=2754  
  Map output bytes=34780  
  Map output materialized bytes=20031  
  Input split bytes=869  
  Combine input records=2754  
  Combine output records=1161  
  Reduce input groups=589  
  Reduce shuffle bytes=20031  
  Reduce input records=1161  
  Reduce output records=589  
  Spilled Records=2322  
  Shuffled Maps =8  
  Failed Shuffles=0  
  Merged Map outputs=8  
  GC time elapsed (ms)=387  
  CPU time spent (ms)=0  
  Physical memory (bytes) snapshot=0  
  Virtual memory (bytes) snapshot=0  
  Total committed heap usage (bytes)=1350402048  
Shuffle Errors  
  BAD_ID=0  
  CONNECTION=0  
  IO_ERROR=0  
  WRONG_LENGTH=0  
  WRONG_MAP=0  
  WRONG_REDUCE=0  
File Input Format Counters  
  Bytes Read=26012  
File Output Format Counters  
  Bytes Written=10079
```

Fig. 10. WordCount result

Check the output(see in Figure 11), press q to quit:

```
less output/*_
```

We have finished building **Single Mode**, and we will start to build **Cluster Mode** of hadoop.

```
hadoop@Master: /usr/local/hadoop
"*" 18
"AS 8
"License"); 8
"alice,bob 18
"kerberos" 1
"simple" 1
'HTTP/' 1
'none' 1
'random' 1
'sasl' 1
'string' 1
'zookeeper' 2
'zookeeper'. 1
(ASF) 1
(Kerberos). 1
(default), 1
(root 1
(specified 1
output/part-r-00000 (file 1 of 2)
```

Fig. 11. WordCount output

## V. CLUSTER MODE CONSTRUCTION

### A. Preparations for constructing cluster mode.

According to the configuration of Hadoop on a single machine, install Hadoop on the other computers in the cluster, including SSH and java environment.

We can check the IP address using(see in Figure 12, the IP address is the first *inet address*. Thus, the IP address of the machine is 10.187.114.78):

```
ifconfig -a
```

```
hadoop@Master:/usr/local/hadoop$ ifconfig -a
eth0      Link encap:Ethernet  HWaddr 00:0c:29:89:14:30
          inet addr:10.187.114.78  Bcast:10.187.255.255  Mask:255.255.0.0
          inet6 addr: fe80::20c:29ff:fe89:1430/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:175776 errors:0 dropped:0 overruns:0 frame:0
          TX packets:439796 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:15808922 (15.8 MB)  TX bytes:923990440 (923.9 MB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:1110 errors:0 dropped:0 overruns:0 frame:0
          TX packets:1110 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:193072 (193.0 KB)  TX bytes:193072 (193.0 KB)
```

Fig. 12. Check IP Address

The addresses of the three machines in the cluster are:

```
10.187.114.78 Master
10.187.114.76 Slave1
10.187.244.157 Slave2
```

Names of your computers may not look like above, then how to modify the names of the machines? Execute the following command on each machine:

```
sudo gedit /etc/hostname
```

Modify the name as *Master*, *Slave1* and *Slave2*, respectively. (One name for one machine)

Then modify the *hosts* file. Correct result is in Figure 13:

```
sudo vim /etc/hosts
```




Fig. 13. Modify the *hosts* file

All the machines should modify the *hosts* file, to ensure that all the *hosts* files are the same.

After the modification, we can use *ping* command to connect to the other two nodes, if the results are similar to Figure 14, then the configuration is done.

```
hadoop@Master:~$ ping Slave1
PING Slave1 (10.187.114.76) 56(84) bytes of data.
64 bytes from Slave1 (10.187.114.76): icmp_seq=1 ttl=64 time=14.5 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=2 ttl=64 time=32.5 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=3 ttl=64 time=51.0 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=4 ttl=64 time=61.8 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=5 ttl=64 time=76.6 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=6 ttl=64 time=26.0 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=7 ttl=64 time=3.79 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=8 ttl=64 time=10.5 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=9 ttl=64 time=99.5 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=10 ttl=64 time=108 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=11 ttl=64 time=31.1 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=12 ttl=64 time=48.0 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=13 ttl=64 time=9.77 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=14 ttl=64 time=92.5 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=15 ttl=64 time=112 ms
64 bytes from Slave1 (10.187.114.76): icmp_seq=16 ttl=64 time=33.0 ms
^C
--- Slave1 ping statistics ---
16 packets transmitted, 16 received, 0% packet loss, time 15039ms
rtt min/avg/max/mdev = 3.797/50.780/112.602/35.846 ms
hadoop@Master:~$
```

Fig. 14. Ping command test

### B. SSH password-free login on slaves

Enable Master node to login on a Slave node without password.

Generate the public key of Master:

```
cd ~/.ssh
ssh-keygen -t rsa
```

Master should be able to *ssh* itself without password, execute:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

Use *ssh* to check the result.(result is in Figure 15):

```
ssh Master
```

Transmit the public key to the two slaves, and password is required in the first time connection:

```
scp /home/hadoop/.ssh/id_rsa.pub hadoop@Slave1: /home/hadoop/
scp /home/hadoop/.ssh/id_rsa.pub hadoop@Slave2: /home/hadoop/
```

On Slave1 and Slave2, save the public key to the target location:

```
cat ~/id_rsa.pub >> ~/.ssh/authorized_keys
```



```
hadoop@Master:~$ ssh Master
Welcome to Ubuntu 14.04.1 LTS (GNU/Linux 3.13.0-39-generic x86_64)

 * Documentation:  https://help.ubuntu.com/

312 packages can be updated.
159 updates are security updates.

Last login: Wed Mar 11 19:07:03 2015 from master
$
```

Fig. 15. SSH command test

Then we can use Master to login on Slave1 and Slave2 without password:

```
ssh Slave1
ssh Slave2
```

### C. Cluster environment configuration

On Master, modify *slaves* file as Figure 16:

```
cd /usr/local/hadoop/etc/hadoop/
sudo gedit slaves
```




Fig. 16. *slaves* file modification

On Master, modify *core-site.xml* file as Figure 17:

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://Master:9000</value>
  </property>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>file:/usr/local/hadoop/tmp</value>
    <description>Abase for other temporary directories.</description>
  </property>
</configuration>
```

Fig. 17. *core-site.xml* file modification

On Master, modify *hdfs-site.xml* file as Figure 18:

The value of *dfs.replication* is the number of slaves.

On Master, copy a template first:

```
cp mapred-site.xml.template mapred-site.xml
```

Then on Master modify *hdfs-site.xml* file as Figure 19:

Finally, on Master, modify *yarn-site.xml* file as Figure 20:

After the configuration of these files, compress the Hadoop file, then send it to slaves.

```
cd /usr/local/
sudo tar -zcf ./hadoop.tar.gz ./hadoop
scp ./hadoop.tar.gz Slave1:/home/hadoop/
scp ./hadoop.tar.gz Slave2:/home/hadoop/
```

```

<configuration>
  <property>
    <name>dfs.namenode.secondary.http-address</name>
    <value>Master:50090</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:/usr/local/hadoop/tmp/dfs/name</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:/usr/local/hadoop/tmp/dfs/data</value>
  </property>
  <property>
    <name>dfs.replication</name>
    <value>2</value>
  </property>
</configuration>

```

Fig. 18. *hdfs-site.xml* file modification

```

<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>

```

Fig. 19. *mapred-site.xml* file modification

On Slave1 and Slave2, execute:

```

sudo tar -zxvf ~/hadoop.tar.gz -C /usr/local/
sudo chown -R hadoop:hadoop /usr/local/hadoop

```

Now, we can set on to run *WordCount* in cluster mode.

## VI. WORDCOUNT TEST IN CLUSTER MODE

***From now on, all the steps are implemented only on Master***

### A. Format HDFS

At the first time of Hadoop operation, we need to format:

```

cd /usr/local/hadoop/
bin/hdfs namenode -format

```

### B. Start HDFS

Then we start hadoop, like in Figure 21:

```

sbin/start-dfs.sh
sbin/start-yarn.sh

```

```

<configuration>
  <property>
    <name>yarn.resourcemanager.hostname</name>
    <value>Master</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>

```

Fig. 20. *yarn-site.xml* file modification

```

hadoop@Master:/usr/local/hadoop$ sbin/start-dfs.sh
15/03/11 10:47:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [Master]
Master: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hadoop-namenode-Master.out
Slave2: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hadoop-datanode-Slave2.out
Slave1: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hadoop-datanode-Slave1.out
Starting secondary namenodes [Master]
Master: starting secondarynamenode, logging to /usr/local/hadoop/logs/hadoop-hadoop-secondarynamenode-Master.out
15/03/11 10:47:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hadoop@Master:/usr/local/hadoop$ sbin/start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /usr/local/hadoop/logs/yarn-hadoop-resourcemanager-Master.out
Slave1: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hadoop-nodemanager-Slave1.out
Slave2: starting nodemanager, logging to /usr/local/hadoop/logs/yarn-hadoop-nodemanager-Slave2.out

```

Fig. 21. Start hadoop in cluster mode

To check processes on each node, we can use (see in Figure 22):

```
jps
```

On *Master*, we should see **NameNode**, **SecondaryNameNode** and **ResourceManager**. On two *Slaves*, we should see **DataNode** and **NodeManager**.

```

hadoop@Master:/usr/local/hadoop$ jps
5929 NameNode
6541 Jps
6138 SecondaryNameNode
6276 ResourceManager

```

Fig. 22. Using jps to check running hadoop processes

To check the status of nodes, we can use (see in Figure 23):

```
bin/hdfs dfsadmin -report
```

We can also check online: <http://master:50070/> (see in Figure 24)

### C. Execute WordCount

Build a new folder in hdfs, then copy everything in `etc/hadoop` into *input*:

```

bin/hdfs dfs -mkdir /user
bin/hdfs dfs -mkdir /user/hadoop
bin/hdfs dfs -put etc/hadoop input

```

Then Run WordCount:

```

bin/hadoop jar share/hadoop/mapreduce/sources/hadoop-mapreduce-examples-2.6.0-sources.jar org.apache.hadoop.examples.WordCount input output

```

We can both check the process using command line and using a webpage: <http://master:8088/cluster> (Figure 25, Figure 26, Figure 27 indicates the beginning, middle and finishing of the process correspondingly). The process result observed using command line is shown in Figure 28. After the execution of *WordCount*, we may observe in Figure 29 that the used storage has changed.

We can check the output result (see in Figure 30):

```
bin/hdfs dfs -cat output/* | less
```

Press `q` to quit.

```

-----
Live datanodes (2):

Name: 10.187.114.76:50010 (Slave1)
Hostname: Slave1
Decommission Status : Normal
Configured Capacity: 14662287360 (13.66 GB)
DFS Used: 24576 (24 KB)
Non DFS Used: 8506228736 (7.92 GB)
DFS Remaining: 6156034048 (5.73 GB)
DFS Used%: 0.00%
DFS Remaining%: 41.99%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Wed Mar 11 19:58:54 CST 2015


Name: 10.187.244.157:50010 (Slave2)
Hostname: Slave2
Decommission Status : Normal
Configured Capacity: 14662287360 (13.66 GB)
DFS Used: 24576 (24 KB)
Non DFS Used: 7711236096 (7.18 GB)
DFS Remaining: 6951026688 (6.47 GB)
DFS Used%: 0.00%
DFS Remaining%: 47.41%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Wed Mar 11 19:58:54 CST 2015

```

Fig. 23. Check the status of Datanodes

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -										
Datanode Information										
In operation										
Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
Slave1 (10.187.114.76:50010)	0	In Service	13.66 GB	104.8 KB	7.92 GB	5.73 GB	35	104.8 KB (0%)	0	2.6.0
Slave2 (10.187.244.157:50010)	1	In Service	13.66 GB	104.8 KB	7.18 GB	6.47 GB	35	104.8 KB (0%)	0	2.6.0

Fig. 24. Check Datanodes online before executing *WordCount*

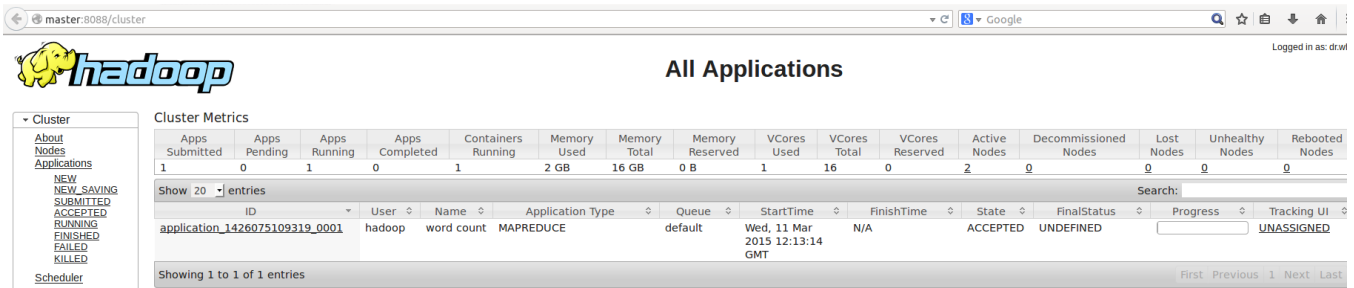


Fig. 25. The beginning of *WordCount*

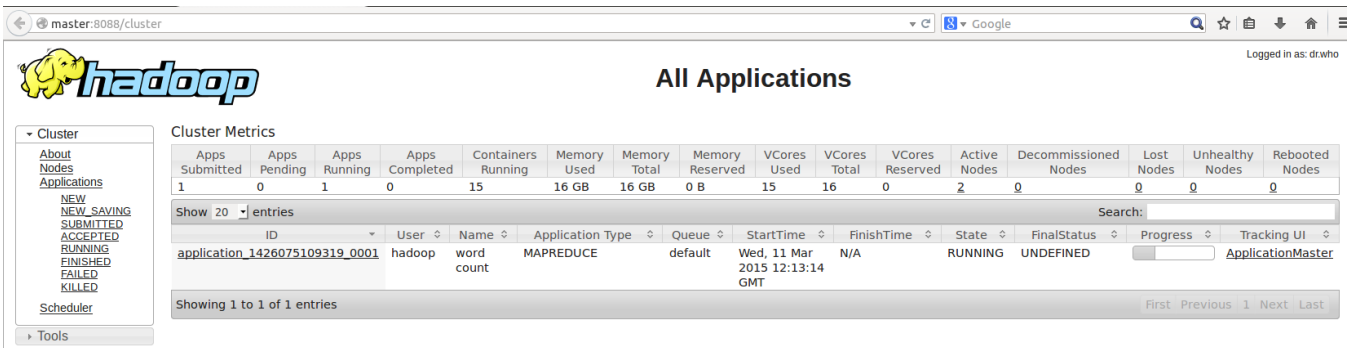


Fig. 26. During the process of *WordCount*

#### D. Stop HDFS

Finally, do not forget to stop hadoop, or strange troubles may occur.(see in Figure 31)

```
sbin/stop-dfs.sh
sbin/stop-yarn.sh
```

OK, that's all the steps to build and run hadoop in a cluster.

### VII. CONCLUSION & ADDITIONAL FINDING

#### A. Conclusion

From this project, we have known the basic steps to construct hadoop in both single mode and cluster mode. Furthermore, we have learned from the project that Hadoop Distributed File System(HDFS) is in a master/slave structure. There is only a single NameNode on one HDFS, which runs on Master. Master is responsible for managing the namespaces of HDFS and the visit from clients. Also, HDFS has many DataNodes, usually one DataNode per slave machine, which is used for storing the data.

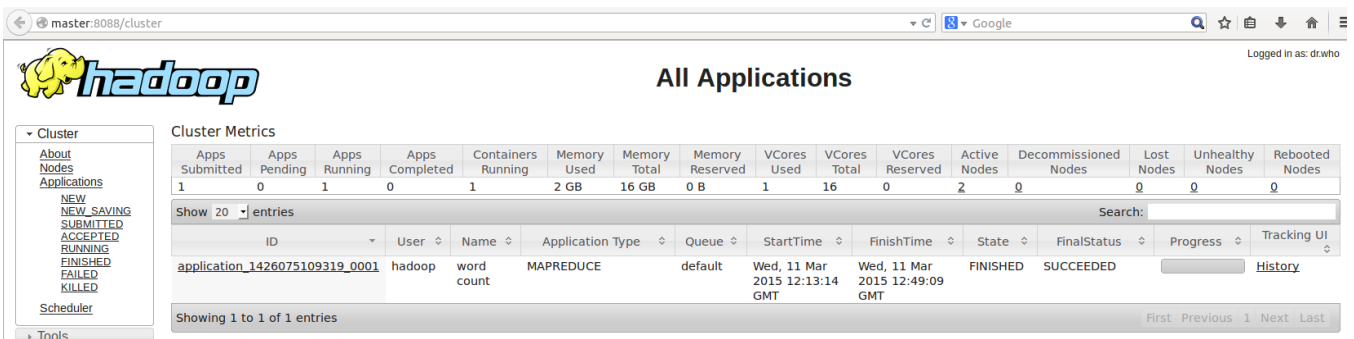


Fig. 27. Finishing the task *WordCount*

```

15/03/11 20:48:02 INFO mapreduce.Job: map 80% reduce 26%
15/03/11 20:48:03 INFO mapreduce.Job: map 86% reduce 26%
15/03/11 20:48:07 INFO mapreduce.Job: map 94% reduce 26%
15/03/11 20:48:09 INFO mapreduce.Job: map 97% reduce 26%
15/03/11 20:48:10 INFO mapreduce.Job: map 97% reduce 31%
15/03/11 20:48:12 INFO mapreduce.Job: map 100% reduce 31%
15/03/11 20:48:13 INFO mapreduce.Job: map 100% reduce 32%
15/03/11 20:48:19 INFO mapreduce.Job: map 100% reduce 33%
15/03/11 20:48:34 INFO mapreduce.Job: map 100% reduce 67%
15/03/11 20:48:44 INFO mapreduce.Job: map 100% reduce 100%
15/03/11 20:49:11 INFO mapreduce.Job: Job job_1426075109319_0001 completed successfully
15/03/11 20:49:14 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=80347
        FILE: Number of bytes written=3962180
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=85917
        HDFS: Number of bytes written=36760
        HDFS: Number of read operations=108
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Failed map tasks=33
        Killed map tasks=2
        Launched map tasks=70
        Launched reduce tasks=1
        Other local map tasks=35
        Data-local map tasks=35
        Total time spent by all maps in occupied slots (ms)=27610230
        Total time spent by all reduces in occupied slots (ms)=915041
        Total time spent by all map tasks (ms)=27610230
        Total time spent by all reduce tasks (ms)=915041
        Total vcore-seconds taken by all map tasks=27610230
        Total vcore-seconds taken by all reduce tasks=915041
        Total megabyte-seconds taken by all map tasks=28272875520
        Total megabyte-seconds taken by all reduce tasks=937001984
    Map-Reduce Framework
        Map input records=2217
        Map output records=8417
        Map output bytes=111668
        Map output materialized bytes=80551
        Input split bytes=4139
        Combine input records=8417
        Combine output records=4356
        Reduce input groups=1576

```

Fig. 28. Check process of *WordCount* using command line

Hadoop Overview Datanodes Snapshot Startup Progress Utilities -										
Datanode Information										
In operation										
Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
Slave1 (10.187.114.76:50010)	2	In Service	13.66 GB	624.89 KB	7.92 GB	5.73 GB	36	624.89 KB (0%)	0	2.6.0
Slave2 (10.187.244.157:50010)	1	In Service	13.66 GB	452.89 KB	7.18 GB	6.47 GB	36	452.89 KB (0%)	0	2.6.0

Fig. 29. Check Datanodes online after executing *WordCount*

```

!=      3
""      6
"."     4
"$HADOOP_CLASSPATH"    1
"$JAVA_HOME"    2
"$YARN_HEAPSIZE"    1
"$YARN_LOGFILE" 1
"$YARN_LOG_DIR" 1
"$YARN_POLICYFILE"    1
"*"      18
"AS      29
"Error: 1
"License");    29
"alice,bob    18
"console"    1
"dfs"    3
"hadoop.root.logger".    1
"jks".    4
"jvm"    3
"mapred"    3
"rpc"    3
"run"    1
"ugi"    3
"x"    1
"x$JAVA_LIBRARY_PATH"    1
#      380
#!/bin/bash    2
###      4
#*.sink.ganglia.dmax=jvm.metrics.threadsBlocked=70,jvm.metrics.memHeapUsedM=40    1
#*.sink.ganglia.slope=jvm.metrics.gcCount=zero,jvm.metrics.memHeapUsedM=both    1
#*.sink.ganglia.tagsForPrefix.dfs=    1
#*.sink.ganglia.tagsForPrefix.jvm=ProcesName    1
#*.sink.ganglia.tagsForPrefix.mapred=    1
#*.sink.ganglia.tagsForPrefix.rpc=    1
#A    1
#Default    1
#HADOOP_JAVA_PLATFORM_OPTS="-XX:-UsePerfData    1
#Security    1
#The    1
:

```

Fig. 30. *WordCount* output

```

hadoop@Master:/usr/local/hadoop$ sbin/stop-dfs.sh
Stopping namenodes on [Master]
Master: stopping namenode
Slave1: stopping datanode
Slave2: stopping datanode
Stopping secondary namenodes [Master]
Master: stopping secondarynamenode
hadoop@Master:/usr/local/hadoop$ sbin/stop-yarn.sh
stopping yarn daemons
stopping resourcemanager
Slave2: stopping nodemanager
Slave1: stopping nodemanager
Slave2: nodemanager did not stop gracefully after 5 seconds: killing with kill -9
Slave1: nodemanager did not stop gracefully after 5 seconds: killing with kill -9
no proxyserver to stop
hadoop@Master:/usr/local/hadoop$ █

```

Fig. 31. Stop HDFS

Files are divided into blocks on DataNodes, usually 64MB per block. Each file may have 2-3 copies separated on different DataNodes to provide storage reliability.

NameNode maintains the namespace, any changes to the category structure of the file system are stored in NameNode. The file system organization is very similar to that of a linux file system. You can create, delete, move and rename files and categories via NameNode. Pay attention that NameNode is only responsible for metadata and there are no data flow from NameNode to clients.

A job is more like an aim we would like to implement using hadoop, while tasks are the subparts of a job. Usually, hadoop divides job into two main kinds of tasks: map tasks and reduce tasks. JobTracker and TaskTracker are two kinds of nodes that control the running of job.

### *B. Additional Finding*

When running *WordCount* in both single mode and cluster mode, we find an interesting phenomenon: it cost several seconds in single mode while cost several minutes in cluster mode. So why the job run much slower using more machines?

This may be caused by the communication cost between seperate machines. Remember that our input data is less than 1MB, so the communication cost far exceeds the computation cost and it turns out that our job runs much faster on a single machine than a distributed cluster.

Hadoop cluster is only suitable for very large number of data(usually TB or PB). When you have to operate on small datasets, using other programming languages and run it on a single machine would be a much better choice.

## REFERENCES

- [1] H. Kopka and P. W. Daly, *A Guide to L<sup>A</sup>T<sub>E</sub>X*, 3rd ed. Harlow, England: Addison-Wesley, 1999.
- [2] *Install Hadoop2.4.0 On Ubuntu14.04 (Single Mode)*, <http://www.cnblogs.com/kinglau/p/3794433.html>
- [3] *Guidance of Installation and Configuration of Hadoop in a cluster: Hadoop2.4.1/Ubuntu 14.04*, <http://www.powerxing.com/install-hadoop-cluster-2-4-1/>